

# Laptop\_data\_Regression

July 2, 2024

## 1 Laptop Price Prediction

**Objective:** To develop a predictive model for estimating the price of laptops based on their specifications. By leveraging various supervised machine learning algorithms, we aim to identify the most accurate model for price prediction, providing valuable insights for both consumers and manufacturers in the laptop market.

**About the Dataset:** The dataset laptop\_data.csv contains the following columns:

- Unnamed: 0: An index column.
- Company: The brand or manufacturer of the laptop (e.g., Dell, Apple, HP).
- TypeName: The type or category of the laptop (e.g., Ultrabook, Gaming, Notebook).
- Inches: The screen size of the laptop in inches.
- ScreenResolution: The resolution of the laptop screen (e.g., 1920x1080, 1366x768).
- Cpu: The central processing unit (CPU) model and speed of the laptop.
- Ram: The amount of random access memory (RAM) in the laptop, typically measured in GB.
- Memory: The storage capacity and type of the laptop (e.g., 256GB SSD, 1TB HDD).
- Gpu: The graphics processing unit (GPU) model in the laptop.
- OpSys: The operating system installed on the laptop (e.g., Windows, macOS, Linux).
- Weight: The weight of the laptop in kilograms.
- Price: The price of the laptop, which is the target variable for prediction.

**Significance:**

- Consumer Decision-Making: By predicting laptop prices based on specifications, consumers can make more informed purchasing decisions and find the best value for their money.
- Market Analysis: Manufacturers and retailers can use the model to analyze market trends, optimize pricing strategies, and understand the impact of various features on the overall price.
- Product Development: Insights gained from the model can guide manufacturers in designing laptops that meet consumer demands while maintaining competitive pricing.

By implementing and comparing various supervised machine learning algorithms, including Linear Regression, Ridge Regression, Lasso Regression, KNeighborsRegressor, DecisionTreeRegressor, SupportVectorRegressor, RandomForestRegressor, and Adaboost Regressor, we aim to identify the most effective model for accurate laptop price prediction.

### 1)Import required libraries

```
[3]: import pandas as pd
import numpy as np
```

```

import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder
from sklearn.metrics import r2_score, mean_absolute_error
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, AdaBoostRegressor
from sklearn.svm import SVR
import pickle
import warnings
warnings.filterwarnings('ignore')

```

## 2 Exploratory Data Analysis

### 2) Read the dataset

```

[91]: df=pd.read_csv('/content/drive/MyDrive/Datasets/laptop_data.csv')
df

```

```

[91]:
    Unnamed: 0  Company      TypeName  Inches  \
0           0   Apple      Ultrabook   13.3
1           1   Apple      Ultrabook   13.3
2           2     HP       Notebook   15.6
3           3   Apple      Ultrabook   15.4
4           4   Apple      Ultrabook   13.3
...         ...      ...
1298        1298  Lenovo  2 in 1 Convertible  14.0
1299        1299  Lenovo  2 in 1 Convertible  13.3
1300        1300  Lenovo       Notebook   14.0
1301        1301     HP       Notebook   15.6
1302        1302   Asus       Notebook   15.6

    ScreenResolution  \
0      IPS Panel Retina Display 2560x1600
1                               1440x900
2      Full HD 1920x1080
3      IPS Panel Retina Display 2880x1800
4      IPS Panel Retina Display 2560x1600
...
1298  IPS Panel Full HD / Touchscreen 1920x1080
1299  IPS Panel Quad HD+ / Touchscreen 3200x1800

```

```

1300                                1366x768
1301                                1366x768
1302                                1366x768

```

```

                                Cpu   Ram           Memory \
0                Intel Core i5 2.3GHz   8GB           128GB SSD
1                Intel Core i5 1.8GHz   8GB   128GB Flash Storage
2            Intel Core i5 7200U 2.5GHz   8GB           256GB SSD
3                Intel Core i7 2.7GHz  16GB           512GB SSD
4                Intel Core i5 3.1GHz   8GB           256GB SSD
...
1298            Intel Core i7 6500U 2.5GHz   4GB           128GB SSD
1299            Intel Core i7 6500U 2.5GHz  16GB           512GB SSD
1300 Intel Celeron Dual Core N3050 1.6GHz   2GB   64GB Flash Storage
1301            Intel Core i7 6500U 2.5GHz   6GB           1TB HDD
1302 Intel Celeron Dual Core N3050 1.6GHz   4GB           500GB HDD

```

```

                                Gpu       OpSys  Weight      Price
0    Intel Iris Plus Graphics 640      macOS  1.37kg  71378.6832
1      Intel HD Graphics 6000      macOS  1.34kg  47895.5232
2      Intel HD Graphics 620       No OS  1.86kg  30636.0000
3      AMD Radeon Pro 455      macOS  1.83kg  135195.3360
4    Intel Iris Plus Graphics 650      macOS  1.37kg  96095.8080
...
1298            Intel HD Graphics 520 Windows 10   1.8kg  33992.6400
1299            Intel HD Graphics 520 Windows 10   1.3kg  79866.7200
1300            Intel HD Graphics  Windows 10   1.5kg  12201.1200
1301            AMD Radeon R5 M330 Windows 10   2.19kg  40705.9200
1302            Intel HD Graphics  Windows 10   2.2kg  19660.3200

```

[1303 rows x 12 columns]

### 3) Statistical Analysis

```
[92]: df.shape
```

```
[92]: (1303, 12)
```

The dataset contains 1303 rows and 12 columns.

```
[93]: df.describe()
```

```

[93]:      Unnamed: 0      Inches      Price
count  1303.00000  1303.000000  1303.000000
mean    651.00000    15.017191  59870.042910
std     376.28801     1.426304  37243.201786
min       0.00000    10.100000   9270.720000
25%     325.50000    14.000000  31914.720000

```

50%	651.00000	15.600000	52054.560000
75%	976.50000	15.600000	79274.246400
max	1302.00000	18.400000	324954.720000

```
[94]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1303 entries, 0 to 1302
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            1303 non-null   int64
1   Company               1303 non-null   object
2   TypeName              1303 non-null   object
3   Inches               1303 non-null   float64
4   ScreenResolution      1303 non-null   object
5   Cpu                  1303 non-null   object
6   Ram                  1303 non-null   object
7   Memory               1303 non-null   object
8   Gpu                  1303 non-null   object
9   OpSys                1303 non-null   object
10  Weight               1303 non-null   object
11  Price                1303 non-null   float64
dtypes: float64(2), int64(1), object(9)
memory usage: 122.3+ KB
```

```
[95]: # to obtain unique elements in each columns.
df.nunique()
```

```
[95]: Unnamed: 0            1303
      Company             19
      TypeName             6
      Inches             18
      ScreenResolution     40
      Cpu                118
      Ram                 9
      Memory            39
      Gpu               110
      OpSys              9
      Weight            179
      Price             791
      dtype: int64
```

### 3 Data Cleaning

#### 4) Handling Missing Values

```
[96]: # To drop a column
df.drop(['Unnamed: 0'],axis=1,inplace=True)
df.drop_duplicates(inplace=True)
df.reset_index(drop=True,inplace=True)
df.head()
```

```
[96]: Company      TypeName  Inches      ScreenResolution \
0   Apple  Ultrabook    13.3  IPS Panel Retina Display 2560x1600
1   Apple  Ultrabook    13.3                        1440x900
2    HP    Notebook    15.6                        Full HD 1920x1080
3   Apple  Ultrabook    15.4  IPS Panel Retina Display 2880x1800
4   Apple  Ultrabook    13.3  IPS Panel Retina Display 2560x1600

      Cpu      Ram      Memory \
0  Intel Core i5 2.3GHz    8GB    128GB SSD
1  Intel Core i5 1.8GHz    8GB  128GB Flash Storage
2  Intel Core i5 7200U 2.5GHz    8GB    256GB SSD
3  Intel Core i7 2.7GHz   16GB    512GB SSD
4  Intel Core i5 3.1GHz    8GB    256GB SSD

      Gpu  OpSys  Weight      Price
0  Intel Iris Plus Graphics 640  macOS  1.37kg  71378.6832
1    Intel HD Graphics 6000  macOS  1.34kg  47895.5232
2    Intel HD Graphics 620   No OS  1.86kg  30636.0000
3    AMD Radeon Pro 455  macOS  1.83kg  135195.3360
4  Intel Iris Plus Graphics 650  macOS  1.37kg  96095.8080
```

```
[97]: df.shape
```

```
[97]: (1274, 11)
```

```
[98]: # To check the missing values
df.isna().sum()
```

```
[98]: Company      0
      TypeName    0
      Inches      0
      ScreenResolution  0
      Cpu          0
      Ram          0
      Memory       0
      Gpu          0
      OpSys        0
      Weight       0
      Price        0
      dtype: int64
```

There is no missing values in the dataset.

## 5)Datatype Conversion

```
[99]: df.dtypes
```

```
[99]: Company          object
      TypeName         object
      Inches          float64
      ScreenResolution object
      Cpu              object
      Ram              object
      Memory           object
      Gpu              object
      OpSys            object
      Weight           object
      Price            float64
      dtype: object
```

```
[100]: #to remove units of Ram and Weight
df['Ram']=df['Ram'].str.replace('GB','')
df['Ram']=df['Ram'].astype(int)
df['Weight']=df['Weight'].str.replace('kg','')
df['Weight']=df['Weight'].astype(float)
df.head()
```

```
[100]: Company  TypeName  Inches  ScreenResolution \
0   Apple  Ultrabook   13.3  IPS Panel Retina Display 2560x1600
1   Apple  Ultrabook   13.3                1440x900
2    HP    Notebook   15.6                Full HD 1920x1080
3   Apple  Ultrabook   15.4  IPS Panel Retina Display 2880x1800
4   Apple  Ultrabook   13.3  IPS Panel Retina Display 2560x1600

      Cpu  Ram  Memory \
0   Intel Core i5 2.3GHz    8    128GB SSD
1   Intel Core i5 1.8GHz    8  128GB Flash Storage
2  Intel Core i5 7200U 2.5GHz    8    256GB SSD
3   Intel Core i7 2.7GHz   16    512GB SSD
4   Intel Core i5 3.1GHz    8    256GB SSD

      Gpu  OpSys  Weight  Price
0  Intel Iris Plus Graphics 640  macOS    1.37  71378.6832
1   Intel HD Graphics 6000  macOS    1.34  47895.5232
2   Intel HD Graphics 620  No OS    1.86  30636.0000
3   AMD Radeon Pro 455  macOS    1.83  135195.3360
4  Intel Iris Plus Graphics 650  macOS    1.37  96095.8080
```


```
[101]: df.dtypes
```

```
[101]: Company          object
      TypeName         object
      Inches          float64
      ScreenResolution object
      Cpu             object
      Ram             int64
      Memory          object
      Gpu             object
      OpSys           object
      Weight          float64
      Price           float64
      dtype: object
```

## 6)Data Visualization and Feature Engineering

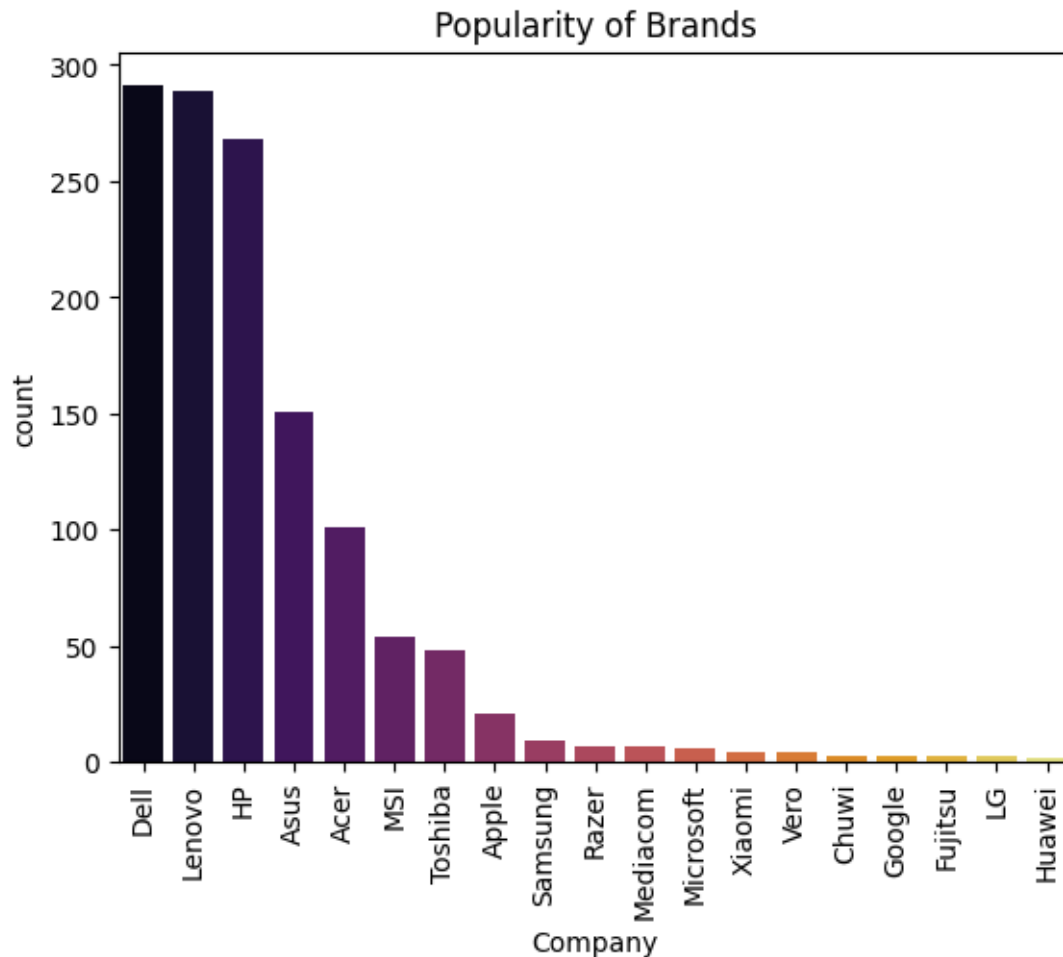
```
[102]: sns.distplot(df['Price'])
      plt.title('Price Distribution')
```

```
[102]: Text(0.5, 1.0, 'Price Distribution')
```



**Insight:** From the distplot, it is clear that the prices in range(0-50000) are more in count.

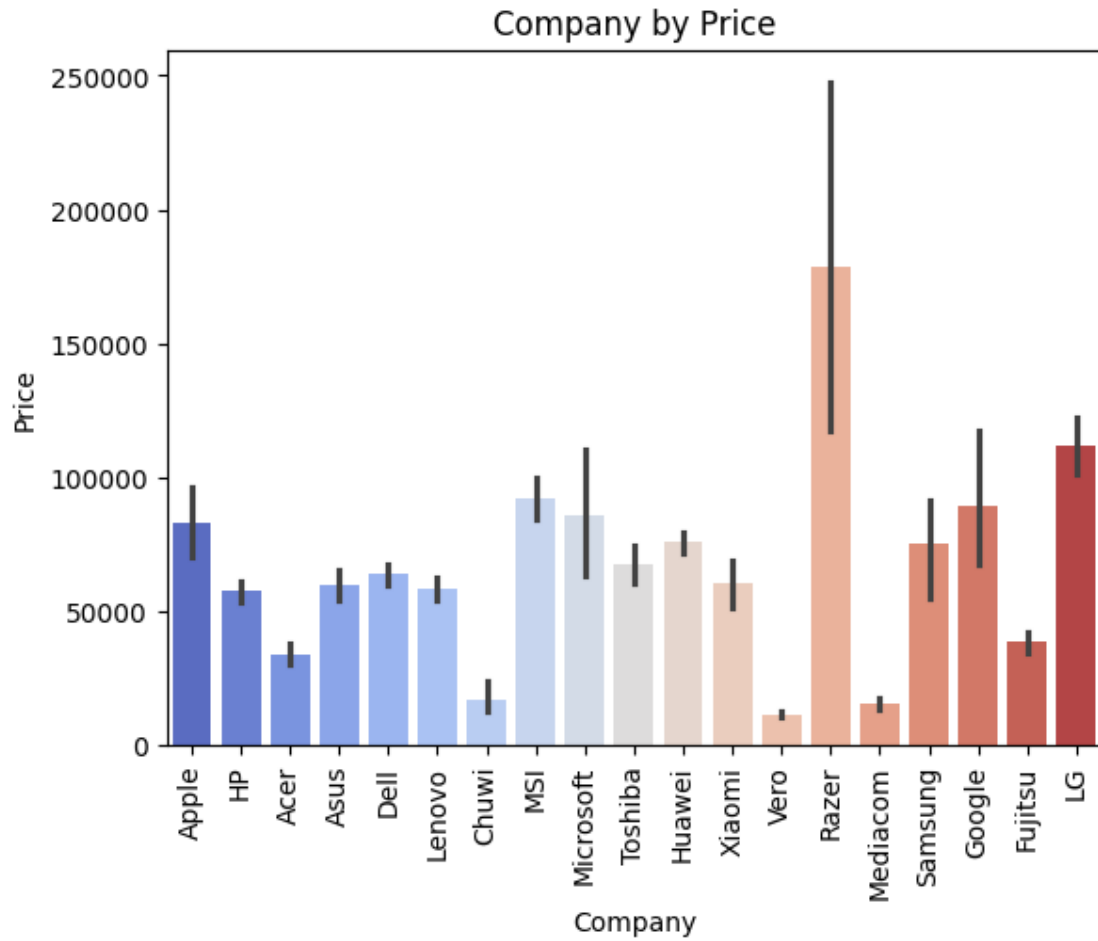
```
[103]: sns.barplot(df['Company'].value_counts(),palette='inferno')
plt.xticks(rotation='vertical')
plt.title('Popularity of Brands')
plt.show()
```



**Insight:** From the barplot, it is evident that majority of the laptops in the dataset are from Dell. Only few number of laptops are of Huawei.


```
[104]: sns.barplot(x=df['Company'],y=df['Price'],palette='coolwarm')
plt.xticks(rotation='vertical')
plt.title('Company by Price')
plt.show()
```





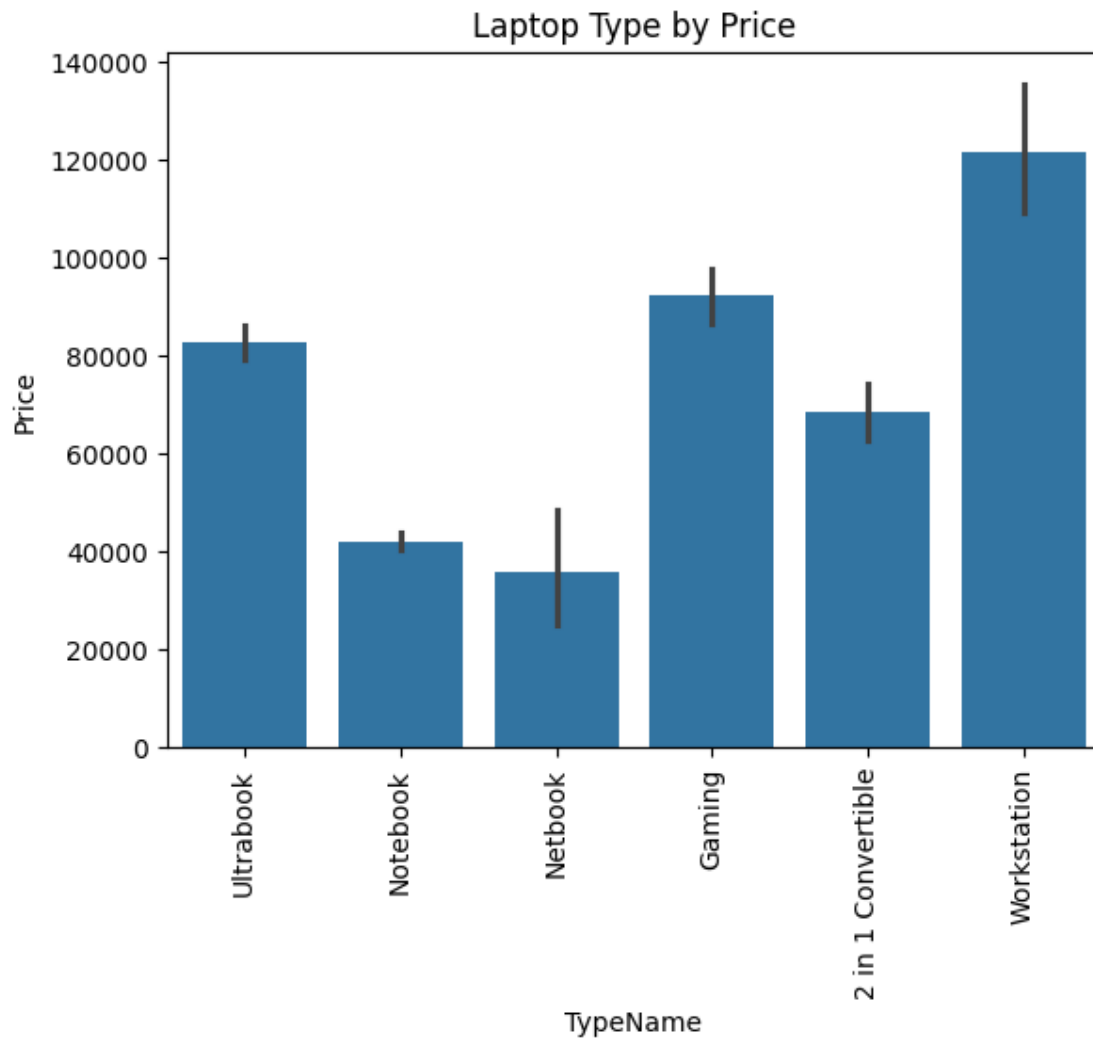
**Insight:** From the barplot, it is clear that Razer laptops are high in price. Whereas Vero brand has low price range.

```
[105]: plt.pie(df['TypeName'].value_counts(),labels=df['TypeName'].value_counts().
        ↪index,
        autopct='%1.1f%%',textprops={'size': 'medium'}, radius=1)
plt.title('Popularity of Laptop Type')
plt.show()
```



**Insight:** From the piechart, it is evident that Notebook type of the laptop are more in count and Netbook type laptops are less in count.


```
[106]: sns.barplot(x=df['TypeName'],y=df['Price'])  
plt.xticks(rotation='vertical')  
plt.title('Laptop Type by Price')  
plt.show()
```



**Insight:** From the barplot, it is clear that Price range is more for Workstation type laptops and less for Netbook type laptops.

```
[107]: sns.scatterplot(x=df['Inches'],y=df['Price'])  
plt.title('Inches by Price')
```

```
[107]: Text(0.5, 1.0, 'Inches by Price')
```



**Insight:** From the scatterplot, it is clear that inches of range(16-18) have high price range.

```
[108]: df['ScreenResolution'].value_counts()
```

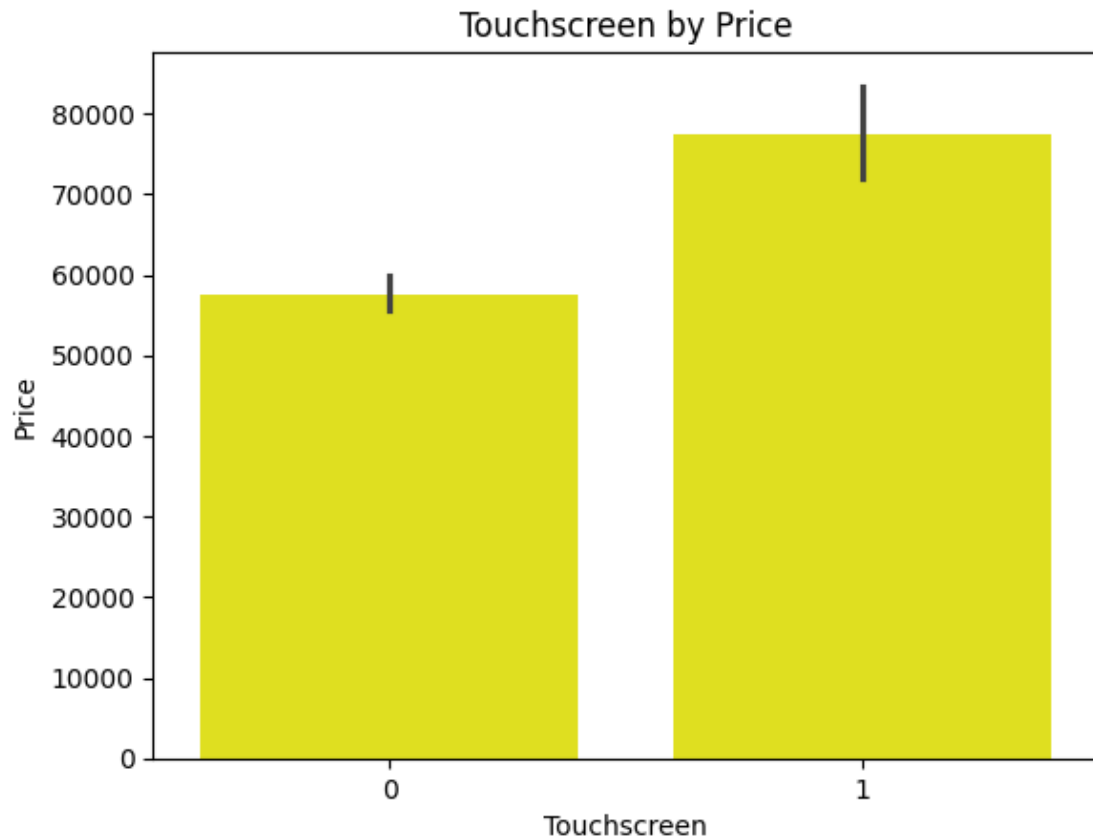
```
[108]: ScreenResolution
Full HD 1920x1080          505
1366x768                  262
IPS Panel Full HD 1920x1080  226
IPS Panel Full HD / Touchscreen 1920x1080  51
Full HD / Touchscreen 1920x1080  47
1600x900                   23
Touchscreen 1366x768       16
Quad HD+ / Touchscreen 3200x1800  15
IPS Panel 4K Ultra HD 3840x2160  12
IPS Panel 4K Ultra HD / Touchscreen 3840x2160  11
4K Ultra HD / Touchscreen 3840x2160  10
4K Ultra HD 3840x2160       7
Touchscreen 2560x1440       7
IPS Panel 1366x768         7
IPS Panel Retina Display 2560x1600  6
IPS Panel Retina Display 2304x1440  6
```

Touchscreen 2256x1504	6
IPS Panel Touchscreen 2560x1440	5
IPS Panel Quad HD+ / Touchscreen 3200x1800	4
IPS Panel Touchscreen 1920x1200	4
1440x900	4
IPS Panel Retina Display 2880x1800	4
IPS Panel 2560x1440	4
2560x1440	3
Quad HD+ 3200x1800	3
1920x1080	3
Touchscreen 2400x1600	3
IPS Panel Quad HD+ 2560x1440	3
IPS Panel Touchscreen 1366x768	3
IPS Panel Touchscreen / 4K Ultra HD 3840x2160	2
IPS Panel Full HD 2160x1440	2
IPS Panel Quad HD+ 3200x1800	2
IPS Panel Retina Display 2736x1824	1
IPS Panel Full HD 1920x1200	1
IPS Panel Full HD 2560x1440	1
IPS Panel Full HD 1366x768	1
Touchscreen / Full HD 1920x1080	1
Touchscreen / Quad HD+ 3200x1800	1
Touchscreen / 4K Ultra HD 3840x2160	1
IPS Panel Touchscreen 2400x1600	1

Name: count, dtype: int64

```
[110]: df['Touchscreen']=df['ScreenResolution'].apply(lambda x:1 if 'Touchscreen' in x_
    ↪else 0)
sns.barplot(x=df['Touchscreen'],y=df['Price'],color='yellow')
plt.title('Touchscreen by Price')
```

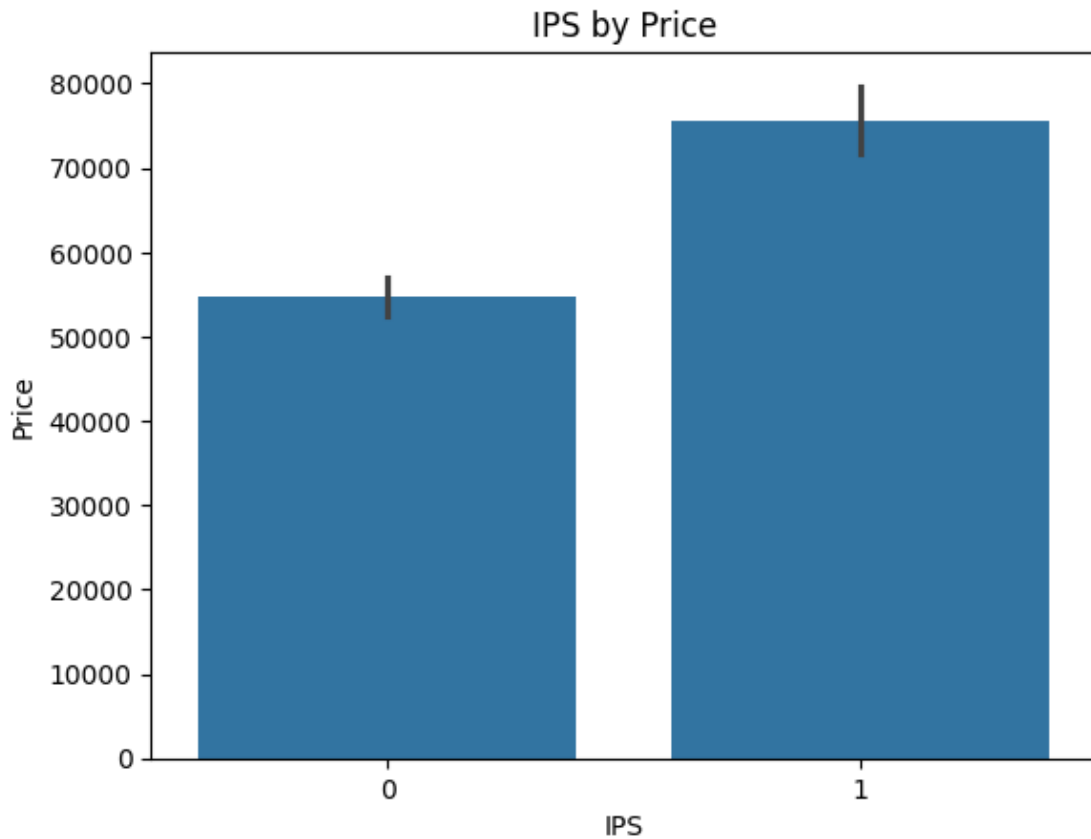
```
[110]: Text(0.5, 1.0, 'Touchscreen by Price')
```



**Insight:** From the barchart, it is clear that laptops with touchscreen have high price range.

```
[35]: df['IPS']=df['ScreenResolution'].apply(lambda x:1 if 'IPS' in x else 0)
sns.barplot(x=df['IPS'],y=df['Price'])
plt.title('IPS by Price')
```

```
[35]: Text(0.5, 1.0, 'IPS by Price')
```



**Insight:** From the barchart, it is evident that laptops with IPS Panel have high price range.

```
[36]: temp=df['ScreenResolution'].str.split('x',n=1,expand=True)
```

```
[37]: df['screen_width']=temp[0]
df['screen_height']=temp[1]
df.head()
```

```
[37]: Company  TypeName  Inches  ScreenResolution \
0   Apple  Ultrabook   13.3  IPS Panel Retina Display 2560x1600
1   Apple  Ultrabook   13.3                1440x900
2    HP    Notebook   15.6                Full HD 1920x1080
3   Apple  Ultrabook   15.4  IPS Panel Retina Display 2880x1800
4   Apple  Ultrabook   13.3  IPS Panel Retina Display 2560x1600

      Cpu  Ram  Memory \
0  Intel Core i5 2.3GHz  8  128GB SSD
1  Intel Core i5 1.8GHz  8  128GB Flash Storage
2  Intel Core i5 7200U 2.5GHz  8  256GB SSD
3  Intel Core i7 2.7GHz 16  512GB SSD
```

			Gpu	OpSys	Weight	Price	Touchscreen	IPS	\
0	Intel	Iris Plus Graphics	640	macOS	1.37	71378.6832	0	1	
1		Intel HD Graphics	6000	macOS	1.34	47895.5232	0	0	
2		Intel HD Graphics	620	No OS	1.86	30636.0000	0	0	
3		AMD Radeon Pro	455	macOS	1.83	135195.3360	0	1	
4	Intel	Iris Plus Graphics	650	macOS	1.37	96095.8080	0	1	

```
[38]: df['screen_width']=df['screen_width'].str.findall(r'(\d+\.\?\d+)').apply(lambda x: x[0])
df.head()
```

		Cpu	Ram	Memory \
0		Intel Core i5 2.3GHz	8	128GB SSD
1		Intel Core i5 1.8GHz	8	128GB Flash Storage
2	Intel	Core i5 7200U 2.5GHz	8	256GB SSD
3		Intel Core i7 2.7GHz	16	512GB SSD
4		Intel Core i5 3.1GHz	8	256GB SSD

			Gpu	OpSys	Weight	Price	Touchscreen	IPS	\
0	Intel	Iris Plus Graphics	640	macOS	1.37	71378.6832	0	1	
1		Intel HD Graphics	6000	macOS	1.34	47895.5232	0	0	
2		Intel HD Graphics	620	No OS	1.86	30636.0000	0	0	
3		AMD Radeon Pro	455	macOS	1.83	135195.3360	0	1	
4	Intel	Iris Plus Graphics	650	macOS	1.37	96095.8080	0	1	

16



```
[39]: df['screen_width']=df['screen_width'].astype(int)
df['screen_height']=df['screen_height'].astype(int)
```

```
[40]: df['PPI']=(((df['screen_width']**2)+(df['screen_height']**2))**0.5/
↳df['Inches']).astype(float)
```

```
[41]: df.dtypes
```

```
[41]: Company          object
      TypeName        object
      Inches          float64
      ScreenResolution object
      Cpu             object
      Ram             int64
      Memory          object
      Gpu             object
      OpSys           object
      Weight          float64
      Price           float64
      Touchscreen     int64
      IPS             int64
      screen_width    int64
      screen_height   int64
      PPI             float64
      dtype: object
```

```
[42]: df.
↳drop(['ScreenResolution','Inches','screen_width','screen_height'],axis=1,inplace=True)
df.head()
```

```
[42]: Company  TypeName          Cpu  Ram          Memory \
0  Apple  Ultrabook      Intel Core i5 2.3GHz    8          128GB SSD
1  Apple  Ultrabook      Intel Core i5 1.8GHz    8  128GB Flash Storage
2    HP   Notebook  Intel Core i5 7200U 2.5GHz    8          256GB SSD
3  Apple  Ultrabook      Intel Core i7 2.7GHz   16          512GB SSD
4  Apple  Ultrabook      Intel Core i5 3.1GHz    8          256GB SSD

      Gpu  OpSys  Weight          Price  Touchscreen  IPS  \
0  Intel Iris Plus Graphics 640  macOS    1.37  71378.6832    0    1
1      Intel HD Graphics 6000  macOS    1.34  47895.5232    0    0
2      Intel HD Graphics 620   No OS    1.86  30636.0000    0    0
3      AMD Radeon Pro 455   macOS    1.83  135195.3360    0    1
4  Intel Iris Plus Graphics 650  macOS    1.37  96095.8080    0    1

      PPI
0  226.983005
1  127.677940
```

```
2 141.211998
3 220.534624
4 226.983005
```

```
[43]: df['Cpu'].value_counts()
```

```
[43]: Cpu
Intel Core i5 7200U 2.5GHz    190
Intel Core i7 7700HQ 2.8GHz   146
Intel Core i7 7500U 2.7GHz    132
Intel Core i7 8550U 1.8GHz     73
Intel Core i5 8250U 1.6GHz     72

...
Intel Core M M3-6Y30 0.9GHz    1
AMD A9-Series 9420 2.9GHz      1
Intel Core i5 2.9GHz           1
AMD A6-Series 7310 2GHz        1
AMD A9-Series 9410 2.9GHz      1
Name: count, Length: 118, dtype: int64
```

```
[44]: df['Cpu Name']=df['Cpu'].apply(lambda x: ' '.join(x.split()[0:3]))
df.head()
```

```
[44]: Company  TypeName          Cpu  Ram          Memory \
0  Apple  Ultrabook      Intel Core i5 2.3GHz    8          128GB SSD
1  Apple  Ultrabook      Intel Core i5 1.8GHz    8  128GB Flash Storage
2  HP     Notebook  Intel Core i5 7200U 2.5GHz    8          256GB SSD
3  Apple  Ultrabook      Intel Core i7 2.7GHz   16          512GB SSD
4  Apple  Ultrabook      Intel Core i5 3.1GHz    8          256GB SSD

          Gpu  OpSys  Weight      Price  Touchscreen  IPS  \
0  Intel Iris Plus Graphics 640  macOS    1.37  71378.6832      0    1
1          Intel HD Graphics 6000  macOS    1.34  47895.5232      0    0
2          Intel HD Graphics 620  No OS    1.86  30636.0000      0    0
3          AMD Radeon Pro 455  macOS    1.83  135195.3360      0    1
4  Intel Iris Plus Graphics 650  macOS    1.37  96095.8080      0    1

          PPI      Cpu Name
0  226.983005  Intel Core i5
1  127.677940  Intel Core i5
2  141.211998  Intel Core i5
3  220.534624  Intel Core i7
4  226.983005  Intel Core i5
```

```
[45]: def Cpu_processor(text):
    if text=='Intel Core i7' or text=='Intel Core i5' or text=='Intel Core i3':
        return text
```

```

else:
    if text.split()[0]=="Intel":
        return 'Other Intel Processor'
    else:
        return 'AMD Processor'

```

```

[46]: df['Cpu Brand']=df['Cpu Name'].apply(Cpu_processor)
df.head()

```

```

[46]: Company   TypeName                Cpu   Ram      Memory \
0   Apple  Ultrabook      Intel Core i5 2.3GHz    8      128GB SSD
1   Apple  Ultrabook      Intel Core i5 1.8GHz    8  128GB Flash Storage
2    HP    Notebook  Intel Core i5 7200U 2.5GHz    8      256GB SSD
3   Apple  Ultrabook      Intel Core i7 2.7GHz   16      512GB SSD
4   Apple  Ultrabook      Intel Core i5 3.1GHz    8      256GB SSD

          Gpu  OpSys  Weight      Price  Touchscreen  IPS \
0  Intel Iris Plus Graphics 640  macOS    1.37  71378.6832         0    1
1          Intel HD Graphics 6000  macOS    1.34  47895.5232         0    0
2          Intel HD Graphics 620   No OS    1.86  30636.0000         0    0
3          AMD Radeon Pro 455  macOS    1.83  135195.3360         0    1
4  Intel Iris Plus Graphics 650  macOS    1.37   96095.8080         0    1


          PPI      Cpu Name      Cpu Brand
0  226.983005  Intel Core i5  Intel Core i5
1  127.677940  Intel Core i5  Intel Core i5
2  141.211998  Intel Core i5  Intel Core i5
3  220.534624  Intel Core i7  Intel Core i7
4  226.983005  Intel Core i5  Intel Core i5

```

```

[48]: sns.barplot(data=df,
                  x='Price',
                  hue='Cpu Brand')
plt.title('Price by CPU Brand')
plt.xlabel('Price')
plt.ylabel('CPU Brand')
plt.show()

```



**Insight:** From the barchart, it is evident that most of the laptops have intel core i7 processor.

```
[49]: df.drop(['Cpu', 'Cpu Name'],axis=1,inplace=True)
df.head()
```

```
[49]: Company  TypeName  Ram      Memory      Gpu \
0  Apple  Ultrabook    8      128GB SSD  Intel Iris Plus Graphics 640
1  Apple  Ultrabook    8  128GB Flash Storage  Intel HD Graphics 6000
2    HP   Notebook    8      256GB SSD  Intel HD Graphics 620
3  Apple  Ultrabook   16      512GB SSD  AMD Radeon Pro 455
4  Apple  Ultrabook    8      256GB SSD  Intel Iris Plus Graphics 650

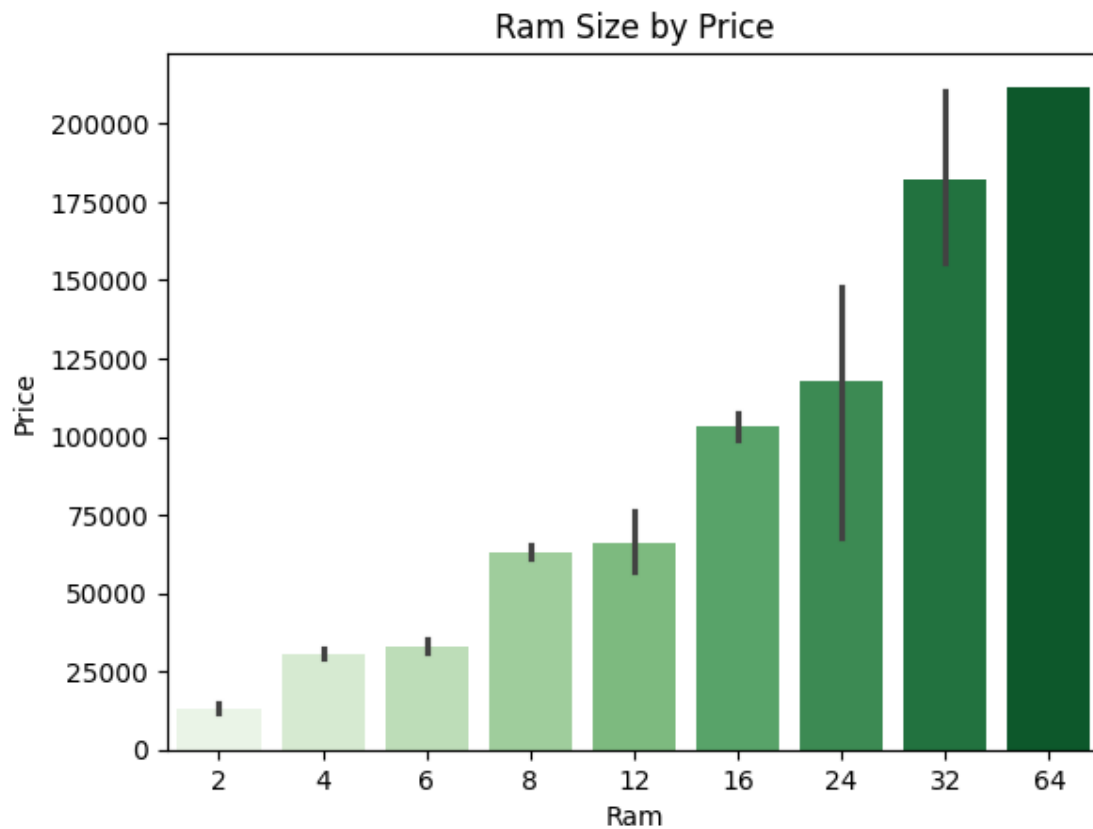
OpSys  Weight      Price  Touchscreen  IPS      PPI      Cpu Brand
0  macOS   1.37  71378.6832           0    1  226.983005  Intel Core i5
1  macOS   1.34  47895.5232           0    0  127.677940  Intel Core i5
2  No OS   1.86  30636.0000           0    0  141.211998  Intel Core i5
3  macOS   1.83  135195.3360           0    1  220.534624  Intel Core i7
4  macOS   1.37  96095.8080           0    1  226.983005  Intel Core i5
```

```
[50]: df['Ram'].value_counts()
```

```
[50]: Ram
      8      613
      4      366
      16     198
      6       35
      12       25
      32       17
      2       16
      24        3
      64        1
      Name: count, dtype: int64
```

```
[52]: sns.barplot(x=df['Ram'],y=df['Price'],palette='Greens')
      plt.title('Ram Size by Price')
```

```
[52]: Text(0.5, 1.0, 'Ram Size by Price')
```



**Insight:** From the barchart, it is clear that most of the laptops have 64GB Ram.

```
[53]: df['Memory'].value_counts()
```

```
[53]: Memory
      256GB SSD          412
      1TB HDD           215
      500GB HDD         123
      512GB SSD         114
      128GB SSD + 1TB HDD  94
      128GB SSD          74
      256GB SSD + 1TB HDD  73
      32GB Flash Storage  36
      2TB HDD            16
      512GB SSD + 1TB HDD  14
      1TB SSD            14
      64GB Flash Storage  13
      256GB SSD + 2TB HDD  10
      256GB Flash Storage   8
      16GB Flash Storage    7
      1.0TB Hybrid         7
      32GB SSD             6
      180GB SSD            5
      128GB Flash Storage   4
      512GB SSD + 2TB HDD   3
      16GB SSD             3
      512GB Flash Storage   2
      1TB SSD + 1TB HDD     2
      256GB SSD + 500GB HDD  2
      128GB SSD + 2TB HDD   2
      256GB SSD + 256GB SSD  2
      512GB SSD + 256GB SSD  1
      512GB SSD + 512GB SSD  1
      64GB Flash Storage + 1TB HDD  1
      1TB HDD + 1TB HDD     1
      32GB HDD             1
      64GB SSD             1
      128GB HDD            1
      240GB SSD            1
      8GB SSD              1
      508GB Hybrid         1
      1.0TB HDD            1
      512GB SSD + 1.0TB Hybrid  1
      256GB SSD + 1.0TB Hybrid  1
      Name: count, dtype: int64
```

```
[54]: df['Memory'] = df['Memory'].astype(str).replace('\.0', '', regex=True)
df["Memory"] = df["Memory"].str.replace('GB', '')
df["Memory"] = df["Memory"].str.replace('TB', '000')
new = df["Memory"].str.split("+", n = 1, expand = True)
```

```
df["first"]= new[0]
df["second"]= new[1]
df.head()
```

```
[54]:
```

	Company	TypeName	Ram	Memory	Gpu	\
0	Apple	Ultrabook	8	128 SSD	Intel Iris Plus Graphics 640	
1	Apple	Ultrabook	8	128 Flash Storage	Intel HD Graphics 6000	
2	HP	Notebook	8	256 SSD	Intel HD Graphics 620	
3	Apple	Ultrabook	16	512 SSD	AMD Radeon Pro 455	
4	Apple	Ultrabook	8	256 SSD	Intel Iris Plus Graphics 650	

	OpSys	Weight	Price	Touchscreen	IPS	PPI	Cpu Brand	\
0	macOS	1.37	71378.6832	0	1	226.983005	Intel Core i5	
1	macOS	1.34	47895.5232	0	0	127.677940	Intel Core i5	
2	No OS	1.86	30636.0000	0	0	141.211998	Intel Core i5	
3	macOS	1.83	135195.3360	0	1	220.534624	Intel Core i7	
4	macOS	1.37	96095.8080	0	1	226.983005	Intel Core i5	

	first	second
0	128 SSD	None
1	128 Flash Storage	None
2	256 SSD	None
3	512 SSD	None
4	256 SSD	None

```
[55]: df['first'].value_counts()
```

```
[55]: first
256 SSD          412
1000 HDD          216
500 HDD           123
512 SSD           114
128 SSD            96
256 SSD            88
128 SSD            74
32 Flash Storage   36
512 SSD            20
2000 HDD           16
1000 SSD           14
64 Flash Storage   13
256 Flash Storage   8
16 Flash Storage    7
1000 Hybrid         7
32 SSD              6
180 SSD             5
128 Flash Storage   4
16 SSD              3
```

```

1000 SSD                2
512 Flash Storage       2
64 SSD                  1
64 Flash Storage        1
1000 HDD                 1
32 HDD                   1
128 HDD                  1
240 SSD                  1
8 SSD                    1
508 Hybrid               1
Name: count, dtype: int64

```

```
[56]: df['second'].value_counts()
```

```

[56]: second
      1000 HDD          185
      2000 HDD           15
      256 SSD            3
      500 HDD            2
      1000 Hybrid         2
      512 SSD             1
Name: count, dtype: int64

```

```

[57]: df["Layer1HDD"] = df["first"].apply(lambda x: 1 if "HDD" in x else 0)
df["Layer1SSD"] = df["first"].apply(lambda x: 1 if "SSD" in x else 0)
df["Layer1Hybrid"] = df["first"].apply(lambda x: 1 if "Hybrid" in x else 0)
df["Layer1Flash_Storage"] = df["first"].apply(lambda x: 1 if "Flash Storage" in
↪x else 0)
df.head()

```

```

[57]: Company  TypeName  Ram          Memory          Gpu  \
0   Apple  Ultrabook    8          128 SSD  Intel Iris Plus Graphics 640
1   Apple  Ultrabook    8  128 Flash Storage  Intel HD Graphics 6000
2    HP    Notebook    8          256 SSD  Intel HD Graphics 620
3   Apple  Ultrabook   16          512 SSD  AMD Radeon Pro 455
4   Apple  Ultrabook    8          256 SSD  Intel Iris Plus Graphics 650

OpSys  Weight      Price  Touchscreen  IPS      PPI      Cpu Brand  \
0  macOS   1.37  71378.6832           0    1  226.983005  Intel Core i5
1  macOS   1.34  47895.5232           0    0  127.677940  Intel Core i5
2  No OS   1.86  30636.0000           0    0  141.211998  Intel Core i5
3  macOS   1.83  135195.3360           0    1  220.534624  Intel Core i7
4  macOS   1.37   96095.8080           0    1  226.983005  Intel Core i5

      first second  Layer1HDD  Layer1SSD  Layer1Hybrid  \
0      128 SSD   None         0          1           0
1  128 Flash Storage  None         0          0           0

```



2	256 SSD	None	0	1	0
3	512 SSD	None	0	1	0
4	256 SSD	None	0	1	0

Layer1Flash_Storage	
0	0
1	1
2	0
3	0
4	0

```
[58]: df['first'] = df['first'].str.findall(r'\d+').apply(lambda x:x[0])
df["second"].fillna("0", inplace = True)
df.head()
```

```
[58]: Company   TypeName  Ram      Memory      Gpu \
0   Apple  Ultrabook    8      128 SSD  Intel Iris Plus Graphics 640
1   Apple  Ultrabook    8  128 Flash Storage  Intel HD Graphics 6000
2    HP    Notebook    8      256 SSD    Intel HD Graphics 620
3   Apple  Ultrabook   16      512 SSD      AMD Radeon Pro 455
4   Apple  Ultrabook    8      256 SSD  Intel Iris Plus Graphics 650
```

	OpSys	Weight	Price	Touchscreen	IPS	PPI	Cpu Brand	\
0	macOS	1.37	71378.6832	0	1	226.983005	Intel Core i5	
1	macOS	1.34	47895.5232	0	0	127.677940	Intel Core i5	
2	No OS	1.86	30636.0000	0	0	141.211998	Intel Core i5	
3	macOS	1.83	135195.3360	0	1	220.534624	Intel Core i7	
4	macOS	1.37	96095.8080	0	1	226.983005	Intel Core i5	

	first	second	Layer1HDD	Layer1SSD	Layer1Hybrid	Layer1Flash_Storage
0	128	0	0	1	0	0
1	128	0	0	0	0	1
2	256	0	0	1	0	0
3	512	0	0	1	0	0
4	256	0	0	1	0	0

```
[59]: df["Layer2HDD"] = df["second"].apply(lambda x: 1 if "HDD" in x else 0)
df["Layer2SSD"] = df["second"].apply(lambda x: 1 if "SSD" in x else 0)
df["Layer2Hybrid"] = df["second"].apply(lambda x: 1 if "Hybrid" in x else 0)
df["Layer2Flash_Storage"] = df["second"].apply(lambda x: 1 if "Flash Storage"
↪in x else 0)
df.head()
```

```
[59]: Company   TypeName  Ram      Memory      Gpu \
0   Apple  Ultrabook    8      128 SSD  Intel Iris Plus Graphics 640
1   Apple  Ultrabook    8  128 Flash Storage  Intel HD Graphics 6000
2    HP    Notebook    8      256 SSD    Intel HD Graphics 620
```

3	Apple	Ultrabook	16	512 SSD	AMD Radeon Pro 455
4	Apple	Ultrabook	8	256 SSD	Intel Iris Plus Graphics 650

	OpSys	Weight	Price	Touchscreen	IPS	...	first	second	Layer1HDD	\
0	macOS	1.37	71378.6832	0	1	...	128	0	0	
1	macOS	1.34	47895.5232	0	0	...	128	0	0	
2	No OS	1.86	30636.0000	0	0	...	256	0	0	
3	macOS	1.83	135195.3360	0	1	...	512	0	0	
4	macOS	1.37	96095.8080	0	1	...	256	0	0	

	Layer1SSD	Layer1Hybrid	Layer1Flash_Storage	Layer2HDD	Layer2SSD	\
0	1	0		0	0	
1	0	0		1	0	
2	1	0		0	0	
3	1	0		0	0	
4	1	0		0	0	

	Layer2Hybrid	Layer2Flash_Storage
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

[5 rows x 22 columns]

```
[60]: df['second'] = df['second'].str.findall(r'\d+').apply(lambda x:x[0])
df.head()
```

```
[60]:
```

	Company	TypeName	Ram	Memory	Gpu	\
0	Apple	Ultrabook	8	128 SSD	Intel Iris Plus Graphics 640	
1	Apple	Ultrabook	8	128 Flash Storage	Intel HD Graphics 6000	
2	HP	Notebook	8	256 SSD	Intel HD Graphics 620	
3	Apple	Ultrabook	16	512 SSD	AMD Radeon Pro 455	
4	Apple	Ultrabook	8	256 SSD	Intel Iris Plus Graphics 650	

	OpSys	Weight	Price	Touchscreen	IPS	...	first	second	Layer1HDD	\
0	macOS	1.37	71378.6832	0	1	...	128	0	0	
1	macOS	1.34	47895.5232	0	0	...	128	0	0	
2	No OS	1.86	30636.0000	0	0	...	256	0	0	
3	macOS	1.83	135195.3360	0	1	...	512	0	0	
4	macOS	1.37	96095.8080	0	1	...	256	0	0	

	Layer1SSD	Layer1Hybrid	Layer1Flash_Storage	Layer2HDD	Layer2SSD	\
0	1	0		0	0	
1	0	0		1	0	
2	1	0		0	0	

3	1	0	0	0	0
4	1	0	0	0	0

	Layer2Hybrid	Layer2Flash_Storage
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

[5 rows x 22 columns]

```
[61]: df['first'].value_counts()
```

```
[61]: first
256      508
1000     240
128      175
512      136
500      123
32        43
2000      16
64        15
16        10
180        5
240        1
8          1
508        1
Name: count, dtype: int64
```

```
[62]: df['second'].value_counts()
```

```
[62]: second
0      1066
1000    187
2000     15
256       3
500       2
512       1
Name: count, dtype: int64
```

```
[63]: df["first"] = df["first"].astype(int)
df["second"] = df["second"].astype(int)
```

```
[64]: df["HDD"]=(df["first"]*df["Layer1HDD"]+df["second"]*df["Layer2HDD"])
df["SSD"]=(df["first"]*df["Layer1SSD"]+df["second"]*df["Layer2SSD"])
df["Hybrid"]=(df["first"]*df["Layer1Hybrid"]+df["second"]*df["Layer2Hybrid"])
```

```
df["Flash_Storage"]=(df["first"]*df["Layer1Flash_Storage"]+df["second"]*df["Layer2Flash_Storage"])
df.drop(columns=['Memory','first','second','Layer1HDD','Layer1SSD','Layer1Hybrid',
                'Layer1Flash_Storage','Layer2HDD','Layer2SSD','Layer2Hybrid',
                'Layer2Flash_Storage'],inplace=True)
df.head()
```

```
[64]: Company   TypeName  Ram      Gpu  OpSys  Weight  \
0   Apple  Ultrabook    8  Intel Iris Plus Graphics 640  macOS    1.37
1   Apple  Ultrabook    8      Intel HD Graphics 6000  macOS    1.34
2    HP    Notebook    8      Intel HD Graphics 620  No OS    1.86
3   Apple  Ultrabook   16      AMD Radeon Pro 455  macOS    1.83
4   Apple  Ultrabook    8  Intel Iris Plus Graphics 650  macOS    1.37

      Price  Touchscreen  IPS      PPI      Cpu Brand  HDD  SSD  Hybrid  \
0  71378.6832           0    1  226.983005  Intel Core i5    0  128      0
1  47895.5232           0    0  127.677940  Intel Core i5    0    0      0
2  30636.0000           0    0  141.211998  Intel Core i5    0  256      0
3  135195.3360           0    1  220.534624  Intel Core i7    0  512      0
4   96095.8080           0    1  226.983005  Intel Core i5    0  256      0

      Flash_Storage
0                0
1               128
2                0
3                0
4                0
```

```
[65]: df.corr(numeric_only=True)['Price']
```

```
[65]: Ram          0.740106
      Weight      0.212192
      Price       1.000000
      Touchscreen  0.188631
      IPS         0.250358
      PPI         0.469539
      HDD        -0.098011
      SSD         0.669957
      Hybrid      0.022533
      Flash_Storage -0.037176
      Name: Price, dtype: float64
```

```
[66]: df.drop(['Hybrid','Flash_Storage'],axis=1,inplace=True)
df.head()
```

```
[66]: Company   TypeName  Ram      Gpu  OpSys  Weight \
0   Apple  Ultrabook    8  Intel Iris Plus Graphics 640  macOS    1.37
1   Apple  Ultrabook    8      Intel HD Graphics 6000  macOS    1.34
2     HP   Notebook    8      Intel HD Graphics 620  No OS    1.86
3   Apple  Ultrabook   16      AMD Radeon Pro 455  macOS    1.83
4   Apple  Ultrabook    8  Intel Iris Plus Graphics 650  macOS    1.37
```

```
      Price  Touchscreen  IPS      PPI      Cpu Brand  HDD  SSD
0  71378.6832           0    1  226.983005  Intel Core i5    0  128
1  47895.5232           0    0  127.677940  Intel Core i5    0    0
2  30636.0000           0    0  141.211998  Intel Core i5    0  256
3 135195.3360           0    1  220.534624  Intel Core i7    0  512
4   96095.8080           0    1  226.983005  Intel Core i5    0  256
```

```
[67]: df['Gpu'].value_counts()
```

```
[67]: Gpu
Intel HD Graphics 620      279
Intel HD Graphics 520      181
Intel UHD Graphics 620      68
Nvidia GeForce GTX 1050      66
Nvidia GeForce GTX 1060      48
...
AMD Radeon R5 520          1
AMD Radeon R7              1
Intel HD Graphics 540      1
AMD Radeon 540             1
ARM Mali T860 MP4          1
Name: count, Length: 110, dtype: int64
```


```
[68]: df['Gpu Brand Name']=df['Gpu'].apply(lambda x:x.split()[0])
df=df[df['Gpu Brand Name']!='ARM']
df.drop('Gpu',axis=1,inplace=True)
df.head()
```

```
[68]: Company   TypeName  Ram  OpSys  Weight      Price  Touchscreen  IPS \
0   Apple  Ultrabook    8  macOS    1.37  71378.6832           0    1
1   Apple  Ultrabook    8  macOS    1.34  47895.5232           0    0
2     HP   Notebook    8  No OS    1.86  30636.0000           0    0
3   Apple  Ultrabook   16  macOS    1.83 135195.3360           0    1
4   Apple  Ultrabook    8  macOS    1.37   96095.8080           0    1


      PPI      Cpu Brand  HDD  SSD  Gpu Brand Name
0  226.983005  Intel Core i5    0  128      Intel
1  127.677940  Intel Core i5    0    0      Intel
2  141.211998  Intel Core i5    0  256      Intel
3  220.534624  Intel Core i7    0  512      AMD
```

4 226.983005 Intel Core i5 0 256 Intel

```
[74]: explode = [0.1,0,0]
df['Gpu Brand Name'].value_counts()
top10_price = df['Gpu Brand Name'].value_counts()
plt.pie(top10_price,
        labels=top10_price.index,
        explode=explode,
        autopct='%1.1f%%')
plt.title('Gpu Brand Name')
plt.show()
```



```
[75]: sns.barplot(x=df['Gpu Brand Name'],y=df['Price'],palette='plasma')
plt.xticks(rotation='vertical')
plt.title('GPU Brand Name by Price')
plt.show()
```



**Insight:** From the barchart, it is clear that price range is high for laptops having Nvidia Gpu Processor.

```
[76]: df['OpSys'].value_counts()
```

```
[76]: OpSys
Windows 10      1047
No OS           66
Linux           58
Windows 7       45
Chrome OS       26
macOS           13
Mac OS X        8
Windows 10 S    8
Android         2
Name: count, dtype: int64
```

```
[77]: def cat_os(inp):
        if inp=='Windows 10' or inp=='Windows 7' or inp=='Windows 10 S':
```

```

        return 'Windows'
    elif inp=='macOS' or inp=='Mac OS X':
        return 'Mac'
    else:
        return 'Others/No OS/Linux'

```

```

[78]: df['OS'] = df['OpSys'].apply(cat_os)
df.drop('OpSys',axis=1,inplace=True)
df.head()

```

```

[78]: Company  TypeName  Ram  Weight      Price  Touchscreen  IPS      PPI  \
0   Apple  Ultrabook    8    1.37  71378.6832             0    1  226.983005
1   Apple  Ultrabook    8    1.34  47895.5232             0    0  127.677940
2    HP    Notebook    8    1.86  30636.0000             0    0  141.211998
3   Apple  Ultrabook   16    1.83  135195.3360             0    1  220.534624
4   Apple  Ultrabook    8    1.37   96095.8080             0    1  226.983005

      Cpu Brand  HDD  SSD Gpu Brand Name      OS
0  Intel Core i5    0  128          Intel      Mac
1  Intel Core i5    0    0          Intel      Mac
2  Intel Core i5    0  256          Intel  Others/No OS/Linux
3  Intel Core i7    0  512           AMD      Mac
4  Intel Core i5    0  256          Intel      Mac


```

```

[80]: sns.barplot(x=df['OS'],y=df['Price'],palette='cividis')
plt.xticks(rotation='vertical')
plt.title('Operating System by Price')
plt.show()

```







**Insight:** From the barchart, it is evident that price range is high for laptops having Mac OS.

```
[81]: sns.histplot(df['Weight'],color='purple')  
      plt.title('Weight Distribution')
```

```
[81]: Text(0.5, 1.0, 'Weight Distribution')
```




```
[82]: sns.scatterplot(x=df['Weight'],y=df['Price'])  
plt.xticks(rotation='vertical')  
plt.title('Weight by Price')  
plt.show()
```




```
[88]: sns.histplot(df['Price'],color='orange')  
plt.title('Price Distribution')
```

```
[88]: Text(0.5, 1.0, 'Price Distribution')
```



```
[89]: sns.histplot(np.log(df['Price']),color='orange')  
plt.title('Price Distribution')
```

```
[89]: Text(0.5, 1.0, 'Price Distribution')
```



## 7) Feature Correlation

```
[ ]: df.corr(numeric_only=True)
```

```
[ ]:
```

	Ram	Weight	Price	Touchscreen	IPS	PPI	\
Ram	1.000000	0.389134	0.739996	0.113316	0.202809	0.294927	
Weight	0.389134	1.000000	0.211667	-0.292288	0.023966	-0.319499	
Price	0.739996	0.211667	1.000000	0.190382	0.251514	0.471481	
Touchscreen	0.113316	-0.292288	0.190382	1.000000	0.136973	0.452107	
IPS	0.202809	0.023966	0.251514	0.136973	1.000000	0.288833	
PPI	0.294927	-0.319499	0.471481	0.452107	0.288833	1.000000	
HDD	0.097340	0.510876	-0.098481	-0.205105	-0.090411	-0.290774	
SSD	0.599552	-0.056985	0.669808	0.252142	0.215197	0.499899	


  

	HDD	SSD
Ram	0.097340	0.599552
Weight	0.510876	-0.056985
Price	-0.098481	0.669808
Touchscreen	-0.205105	0.252142
IPS	-0.090411	0.215197

```
PPI          -0.290774  0.499899
HDD           1.000000 -0.400625
SSD          -0.400625  1.000000
```

```
[ ]: sns.heatmap(df.corr(numeric_only=True))
```

```
[ ]: <Axes: >
```



## 8) Splitting data to X and y

```
[ ]: Ind_var=df.drop(columns=['Price'])
      Ind_var
```

```
[ ]:
   Company  TypeName  Ram  Weight  Touchscreen  IPS  PPI  \
0   Apple   Ultrabook   8    1.37           0    1  226.983005
1   Apple   Ultrabook   8    1.34           0    0  127.677940
2    HP     Notebook   8    1.86           0    0  141.211998
3   Apple   Ultrabook  16    1.83           0    1  220.534624
4   Apple   Ultrabook   8    1.37           0    1  226.983005
```

...	...	...	...	...	...	...	...	...
1269	Asus		Notebook	4	2.20		0	0 100.454670
1270	Lenovo	2 in 1	Convertible	4	1.80		1	1 157.350512
1271	Lenovo	2 in 1	Convertible	16	1.30		1	1 276.053530
1272	Lenovo		Notebook	2	1.50		0	0 111.935204
1273	HP		Notebook	6	2.19		0	0 100.454670

	Cpu Brand	HDD	SSD	Gpu Brand Name	OS
0	Intel Core i5	0	128	Intel	Mac
1	Intel Core i5	0	0	Intel	Mac
2	Intel Core i5	0	256	Intel	Others/No OS/Linux
3	Intel Core i7	0	512	AMD	Mac
4	Intel Core i5	0	256	Intel	Mac
...	...	...	...	...	...
1269	Intel Core i7	500	0	Nvidia	Windows
1270	Intel Core i7	0	128	Intel	Windows
1271	Intel Core i7	0	512	Intel	Windows
1272	Other Intel Processor	0	0	Intel	Windows
1273	Intel Core i7	1000	0	AMD	Windows

[1273 rows x 12 columns]

```
[ ]: Dep_var=np.log(df['Price'])
Dep_var
```

```
[ ]: 0      11.175755
      1      10.776777
      2      10.329931
      3      11.814476
      4      11.473101

      ...
1269    10.555257
1270    10.433899
1271    11.288115
1272     9.409283
1273    10.614129
Name: Price, Length: 1273, dtype: float64
```

## 9) Splitting data to train data and test data

```
[ ]: X_train,X_test,y_train,y_test=train_test_split(Ind_var,Dep_var,test_size=0.
↪10,random_state=2)
```

## 4 Building ML models

```
[ ]: step1 = ColumnTransformer(transformers=[('col_tnf', OneHotEncoder(sparse=False, drop='first'), [0,1,7,
lr=LinearRegression()
rg=Ridge(alpha=10, random_state=42)
ls=Lasso(alpha=0.001, random_state=42)
knn=KNeighborsRegressor(n_neighbors=3)
dt=DecisionTreeRegressor(max_depth=8, random_state=42)
sv=SVR(kernel='rbf')
rf=RandomForestRegressor(n_estimators=100, random_state=42, max_samples=0.
    5, max_features=0.75, max_depth=15)
ad=AdaBoostRegressor(n_estimators=15, random_state=42)

step2 = [lr, rg, ls, knn, dt, sv, rf, ad]
for i in step2:
    print('_'*50, i, '_'*50)
    pipe = Pipeline([('step1', step1), ('step2', i)])
    pipe.fit(X_train, y_train)
    y_pred = pipe.predict(X_test)
    print('R2 Score', r2_score(y_test, y_pred))
    print('MAE', mean_absolute_error(y_test, y_pred))
```

----- LinearRegression()

R2 Score 0.7926901156667799

MAE 0.22065476517716653

----- Ridge(alpha=10,  
random\_state=42) -----

R2 Score 0.7927938535467675

MAE 0.2235451409312221

----- Lasso(alpha=0.001,  
random\_state=42) -----

R2 Score 0.7911292878893352

MAE 0.22399649007048428

----- KNeighborsRegressor(n\_neighbors=3)

R2 Score 0.7486634726488367

MAE 0.2154187555553761

----- DecisionTreeRegressor(max\_depth=8, random\_state=42)

R2 Score 0.8019037046093079

MAE 0.2164088620174775

----- SVR()

-----



```
R2 Score 0.6515291902558747
MAE 0.26461314868916874
```

```
-----
RandomForestRegressor(max_depth=15, max_features=0.75, max_samples=0.5,
                      random_state=42)
```

```
-----
R2 Score 0.8525464607117873
MAE 0.1794891533537441
```

```
-----
AdaBoostRegressor(n_estimators=15, random_state=42)
```

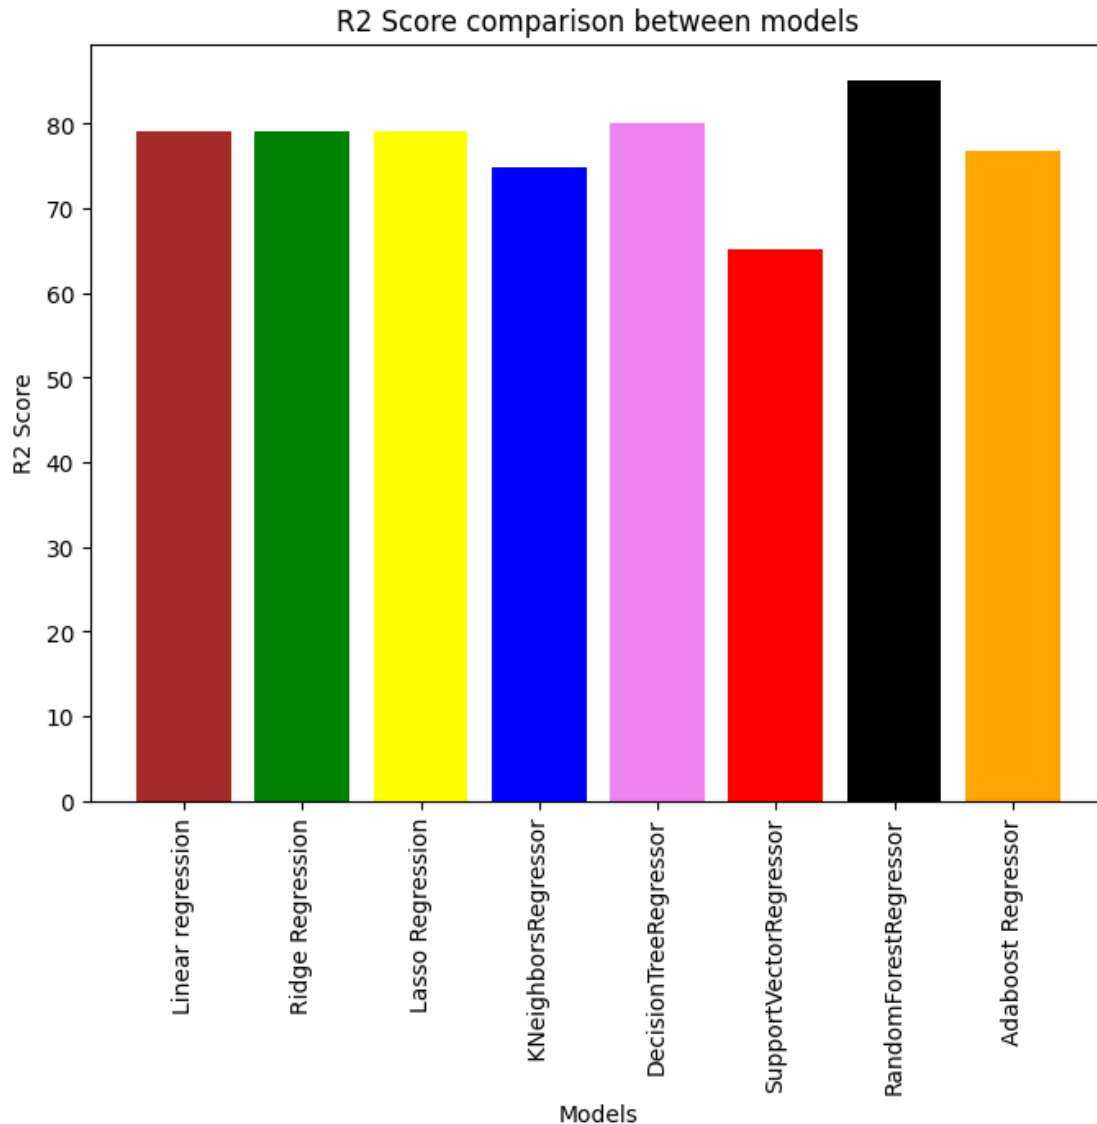
```
-----
R2 Score 0.7688691453495048
MAE 0.2500654165508208
```

```
[ ]: # Accuracy Comparison
Table=pd.DataFrame({'Model':['Linear regression','Ridge Regression','Lasso Regression',
                              'KNeighborsRegressor','DecisionTreeRegressor','SupportVectorRegressor','RandomForestRegressor',
                              'AdaBoostRegressor'],
                    'R2 Score':[0.792,0.792,0.791,0.748,0.801,0.651,0.852,0.768],
                    'MAE':[0.220,0.223,0.223,0.215,0.216,0.264,0.179,0.250]})
Table['R2 Score']=Table['R2 Score']*100
Table
```

```
[ ]:
```

	Model	R2 Score	MAE
0	Linear regression	79.2	0.220
1	Ridge Regression	79.2	0.223
2	Lasso Regression	79.1	0.223
3	KNeighborsRegressor	74.8	0.215
4	DecisionTreeRegressor	80.1	0.216
5	SupportVectorRegressor	65.1	0.264
6	RandomForestRegressor	85.2	0.179
7	Adaboost Regressor	76.8	0.250

```
[ ]: # Plotting Accuracy of Models
plt.figure(figsize=(8,6))
Models=['Linear regression','Ridge Regression','Lasso Regression','KNeighborsRegressor',
        'DecisionTreeRegressor','SupportVectorRegressor','RandomForestRegressor',
        'AdaBoostRegressor']
r2_score=[79.2,79.2,79.1,74.8,80.1,65.1,85.2,76.8]
plt.title('R2 Score comparison between models')
plt.xlabel('Models')
plt.ylabel('R2 Score')
plt.
    bar(Models,r2_score,color=['brown','green','yellow','blue','violet','red','black','orange'])
plt.xticks(rotation='vertical')
plt.show()
```



**Insight:** From the bar chart, it is evident that Random Forest Regressor model has high R2 Score and low MAE.

### Predicting the Laptop Price

```
[ ]: # Selecting the best performing model and training it again with data for prediction
      ↪ prediction
      best_model = RandomForestRegressor(n_estimators=100, random_state=3,
      ↪ max_samples=0.5, max_features=0.75, max_depth=15)
      rf_pipe = Pipeline([('step1', step1), ('step2', best_model)])
      rf_pipe.fit(X_train, y_train)
```

```
[ ]: Pipeline(steps=[('step1',
                        ColumnTransformer(remainder='passthrough',
                                           transformers=[('col_tnf',
                                                           OneHotEncoder(drop='first',
                                                                           sparse=False),
                                                           [0, 1, 7, 10, 11]))]),
                    ('step2',
                     RandomForestRegressor(max_depth=15, max_features=0.75,
                                           max_samples=0.5, random_state=3))])
```

```
[ ]: sample = pd.DataFrame([['Apple', 'Ultrabook', '8', '1.37', '0', '1', '226.
↪983005', 'Intel Core i5', '0', '128', 'Intel', 'Mac']],
                           columns=['Company', 'TypeName', 'Ram', 'Weight',
↪ 'Touchscreen', 'IPS', 'PPI', 'Cpu Brand', 'HDD', 'SSD', 'Gpu Brand Name',
↪ 'OS'])
y_new = rf_pipe.predict(sample)
print('Predicted Price:', y_new)
```

Predicted Price: [11.13052688]

```
[ ]: Price=np.exp(y_new)
print("The predicted price of this configuration is ",Price)
```

The predicted price of this configuration is [68222.3074989]

### Exporting the model

```
[ ]: pickle.dump(rf_pipe,open('rf_pipe.pkl','wb'))
pickle.dump(df,open('df.pkl','wb'))
```

### Deploying the model using Streamlit

```
[ ]: !pip install streamlit -q
```

```
8.6/8.6 MB
19.7 MB/s eta 0:00:00
207.3/207.3
kB 17.6 MB/s eta 0:00:00
6.9/6.9 MB
43.2 MB/s eta 0:00:00
83.0/83.0 kB
8.1 MB/s eta 0:00:00
62.7/62.7 kB
6.7 MB/s eta 0:00:00
```

```
[ ]: %%writefile app.py
import streamlit as st
```

```

from PIL import Image
import pickle
import numpy as np

# Custom CSS to set background color and font color
page_bg_img = '''
<style>
body {
    background-color: #000000;
    color: #ffffff;
}
/* Style for select boxes */
div[data-baseweb="select"] > div:first-child, .st-ct {
    background-color: #000000 !important;
    color: #ffffff !important;
}

/* Style for number input */
.stNumberInput input {
    background-color: #000000;
    color: #ffffff;
}
</style>
'''
st.markdown(page_bg_img, unsafe_allow_html=True)

# import the model
pipe = pickle.load(open('/content/drive/MyDrive/saved models/rf_pipe.pkl','rb'))
df = pickle.load(open('/content/drive/MyDrive/saved models/df.pkl','rb'))

st.title("Laptop Price Predictor")
st.subheader("Welcome! Here you can find out the price of your desired Laptop.")
image=Image.open('/content/drive/MyDrive/laptop_img.jpeg')
st.image(image, width=800)

st.subheader("Fill in the fields")

# brand
company = st.selectbox('Brand',df['Company'].unique())

# type of laptop
typename = st.selectbox('Type',df['TypeName'].unique())

# Ram
ram = st.selectbox('RAM(in GB)', [2,4,6,8,12,16,24,32,64])

# weight

```

```

weight = st.number_input('Weight of the Laptop')

# Touchscreen
touchscreen = st.selectbox('Touchscreen', ['No', 'Yes'])

# IPS
ips = st.selectbox('IPS', ['No', 'Yes'])

# screen size
inches = st.number_input('Screen Size')

# resolution
screen_resolution = st.selectbox('Screen_Resolution', ['1920x1080', '1366x768', '1600x900', '3840x2160', '3200x1800', '2880x1800', '2560x1600'])

#cpu
cpu = st.selectbox('CPU', df['Cpu Brand'].unique())

hdd = st.selectbox('HDD(in GB)', [0, 128, 256, 512, 1024, 2048])

ssd = st.selectbox('SSD(in GB)', [0, 8, 128, 256, 512, 1024])

gpu = st.selectbox('GPU', df['Gpu Brand Name'].unique())

os = st.selectbox('OS', df['OS'].unique())

if st.button('Predict Price'):

    if touchscreen == 'Yes':
        touchscreen = 1
    else:
        touchscreen = 0

    if ips == 'Yes':
        ips = 1
    else:
        ips = 0

    screen_width= int(screen_resolution.split('x')[0])
    screen_height= int(screen_resolution.split('x')[1])
    ppi = ((screen_width**2) + (screen_height**2))**0.5/inches

    # create the query array
    query = np.
    ↪array([company,typename,ram,weight,touchscreen,ips,ppi,cpu,hdd,ssd,gpu,os],dtype=object).
    ↪reshape(1, -1)

```

```
# make the prediction
predicted_price = int(np.exp(pipe.predict(query)[0]))
st.title(f"The predicted price of this configuration is {predicted_price}")
```

Writing app.py

```
[ ]: !wget -q -O - ipv4.icanhazip.com
```

34.168.116.153

```
[ ]: ! streamlit run app.py & npx localtunnel --port 8501
```

Collecting usage statistics. To deactivate, set `browser.gatherUsageStats` to `false`.

You can now view your Streamlit app in your browser.

Local URL: <http://localhost:8501>

Network URL: <http://172.28.0.12:8501>

External URL: <http://34.168.116.153:8501>

npx: installed 22 in 7.044s

your url is: <https://common-masks-punch.local.lt>

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but OneHotEncoder was fitted with feature names  
warnings.warn(

Stopping...

^C